


Parallel Ray Tracing

CUDA



Idea:









1. Execute ray simulation in iterations
2. Reassign new reflected and refracted rays each iteration
3. Compute surface colors from last layer and iterate backwards

Constraints:

- Each layer depends on previous layer
- Extremely memory intensive
- Dynamically scaling workload per layer
- Diffuse vs bounced rays not data parallel

Benefits


- Collision checking is data parallel
- Tracing each ray is independent



Conclusions:


- More memory constrained the compute constrained
- Increasing object count less impactful than expected
- Ideal platform for ray tracing implementation

Background



The ray tracing algorithm is a way to render a 3D scene to a 2D image by simulating the mathematical properties of light rays.

Benchmarks



Three Load Factors:

1. Resolution: 640x460, 1280x720, 1920x1080
2. Number of objects: 1 to 201
3. Number of non-diffusing rays: PAN vs SPIN

Naive Implementation

Idea:



For each ray, recursively trace it and all its scattered rays until ray depth or diffuse, then recursively backtrack to compute surface color.

Benefits:

- Simple to understand
- Create a complex and accurate simulation of real light and shadows

Drawbacks:

- Must check every object against every ray
- Recursion is expensive in practice
- Slow compared to rasterization



OpenMP/AVX

Idea:









1. Execute multiple ray simulations in parallel
2. Distribute workload evenly between threads
3. Optimize arithmetic computations with vector intrinsics

Constraints:

- Each layer depends on previous layer
- Nested and recursive structure harder to parallelize with OpenMP
- Using AVX intrinsics with complex data types

Benefits


- Tracing each ray is independent
- Dynamic scheduling prevents large load



Conclusions:

- Increasing object count highly impactful versus CUDA
- Provides general speedup benefits over naive, but at a completely lower scale than CUDA

CUDA



Idea:

1. Execute ray simulation in iterations
2. Reassign new reflected and refracted rays each iteration
3. Compute surface colors from last layer and iterate backwards

Constraints:


- Each layer depends on previous layer
- Extremely memory intensive
- Dynamically scaling workload per layer
- Diffuse vs bounced rays not data parallel

Benefits

- Collision checking is data parallel
- Tracing each ray is independent

Parallel Ray Tracing

Background



The ray tracing algorithm is a way to render a 3D scene to a 2D image by simulating the mathematical properties of light rays.

OpenMP/AVX

Idea:

1. Execute multiple ray simulations in parallel
2. Distribute workload evenly between threads
3. Optimize arithmetic computations with vector intrinsics


Constraints:

- Each layer depends on previous layer
- Nested and recursive structure harder to parallelize with OpenMP
- Using AVX intrinsics with complex data types


Benefits

- Tracing each ray is independent
- Dynamic scheduling prevents large load imbalances


CUDA Naive vs Naive Pan Speedup




CUDA Naive vs Naive Spin Speedup




CUDA Final vs CUDA Naive Pan Speedup



CUDA Final vs CUDA Naive Spin Speedup




Benchmarks




Three Load Factors:

1. Resolution: 640x460, 1280x720, 1920x1080
2. Number of objects: 1 to 201
3. Number of non-diffusing rays: PAN vs SPIN


OpenMP Final Pan Time



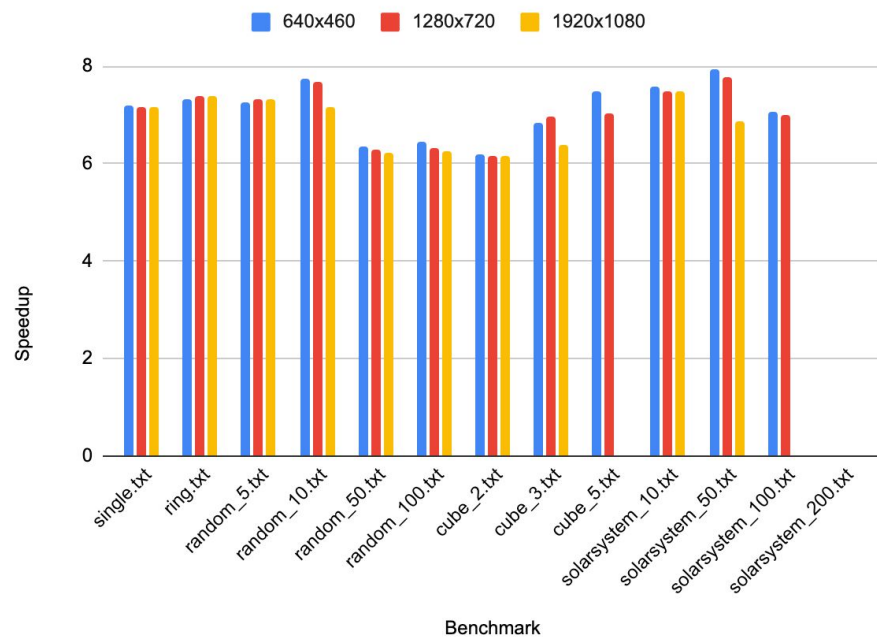
OpenMP Final Spin Time




OpenMP Final Pan vs Initial Speedup




OpenMP Final Spin vs Initial Speedup




CUDA Final Pan Execution Time Breakdown




CUDA Final Spin Execution Time Breakdown



CUDA Final Pan FPS



CUDA Final Spin FPS



Conclusions:

- More memory constrained the compute constrained
- Increasing object count less impactful than expected
- Ideal platform for ray tracing implementation

Naive Implementation

Idea:

For each ray, recursively trace it and all its scattered rays until ray depth or diffuse, then recursively backtrack to compute surface color.

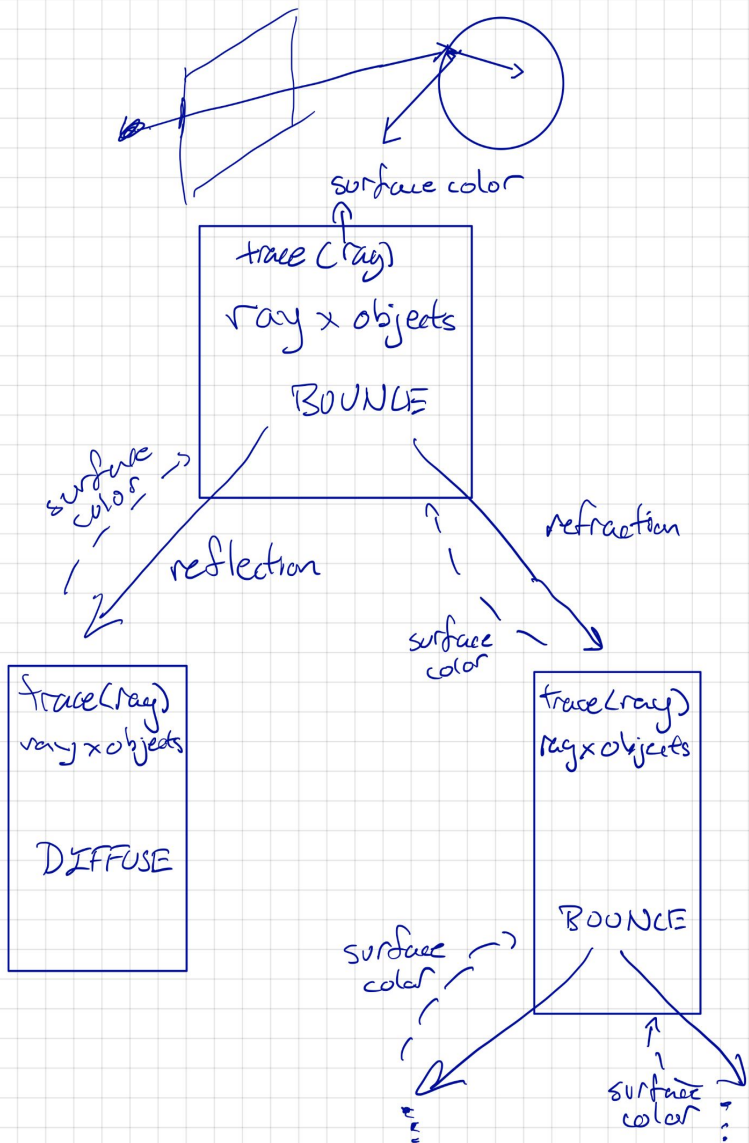

Benefits:

- Simple to understand
- Create a complex and accurate simulation of real light and shadows


Drawbacks:

- Must check every object against every ray
- Recursion is expensive in practice
- Slow compared to rasterization


Naive Pan Time




OpenMP Final Pan FPS




OpenMP Final Spin FPS



OpenMP vs CUDA Pan Speedup



OpenMP vs CUDA Spin Speedup



Conclusions:

- Increasing object count highly impactful versus CUDA
- Provides general speedup benefits over naive, but at a completely lower scale than CUDA